

HARSHIT TRIPATHI

Phone: +919035130434 | Email: ht1992@gmail.com
Location: Gurugram, India | LinkedIn: linkedin.com/in/harshittripathi

PROFESSIONAL SUMMARY

Data Engineer with 10+ years architecting and leading scalable data platforms processing terabytes across AWS, Azure, and GCP. Databricks Certified Data Engineering Associate specializing in real-time streaming, ETL optimization, and data governance with proven impact: 60 percent deployment acceleration, 45 percent cost reduction, and 99.9 percent pipeline reliability. Expert in Databricks, PySpark, Python, Apache Spark, and cloud-native architectures delivering data solutions for 500+ enterprise users.

TECHNICAL SKILLS / CORE COMPETENCIES

Languages & Frameworks: Python, PySpark, SQL, Spark Streaming, Delta Lake, Apache Kafka, Unix Shell Scripting

Big Data & Cloud Platforms: Databricks, Apache Spark, Delta Live Tables (DLT), Unity Catalog, AWS (S3, Redshift, MWAA, Athena, Batch), Azure (ADLS, Data Factory, Synapse), GCP (BigQuery, Cloud Storage)

Databases & Data Warehouses: PostgreSQL, Oracle 11g/12g, MS SQL Server, BigQuery, Amazon Redshift, Azure SQL Data Warehouse

ETL & Data Integration: Databricks Delta Live Tables, Informatica PowerCenter, Informatica Cloud (ICS), CDC pipelines, Real-time streaming, Data pipeline development

DevOps & Orchestration: Apache Airflow, CI/CD pipelines, Harness, GitHub Actions, GitLab CI, Databricks Asset Bundles (DAB), Autosys, Apache Oozie, Kubernetes, Docker, Terraform

Data Governance & Compliance: Unity Catalog, GDPR compliance, Data quality frameworks, Audit monitoring, Data lineage tracking

PROFESSIONAL EXPERIENCE

Publicis Sapient

Senior Associate Data Engineering L2

Gurugram, India
July 2024 – Present

- Led a core team of 6 associates in architecting and deploying a highly scalable Tier 1 platform to drive performance and reliability.
- Architected Platform Data Product on Databricks (AWS/Azure) serving 500+ consumers with sub-second query latency through optimized Delta Lake architecture with Dremio integration
- Engineered automated metadata extraction pipeline for 50+ Unity Catalog and non-Unity workspaces using Account API and Workspace Client API, improving governance visibility by 40 percent and eliminating 25 hours of monthly manual audits
- Built end-to-end CI/CD pipelines with Harness and Apache Airflow orchestrating Databricks Asset Bundle (DAB) workflows, reducing deployment cycles by 99 percent from 4 hours to 2 minutes through automation via Copier templating
- Designed Hive Sync Unity Catalog framework onboarding 15,000+ external Hive tables to Unity Catalog without manual monitoring requirements
- Developed observability dashboards analyzing Unity Catalog system audit tables, monitoring 1 million+ daily queries for compliance, performance optimization, and cost tracking
- Have used Azure DevOps to maintain boards and pipelines for continuous deployments. This was done to deploy data bricks job/workflows and certain assets.

Enquero (A Genpact Company)

Manager - Data Engineering

Gurugram, India
March 2023 – July 2024

- Led a team of 7 Big Data ETL initiative on AWS Batch, orchestrating 15+ production pipelines processing 1.5TB daily retail transactions with 99.9 percent accuracy and zero data loss through robust error handling
- Implemented Slowly Changing Dimension (SCD) Type 1 transformation framework processing 50 million+ daily records from multi-source streams, achieving 35 percent faster processing through Delta Merge operations and efficient partitioning

- Designed comprehensive data quality validation strategy across 200+ Delta tables, achieving 99.8 percent data integrity through automated reconciliation between source and transformed datasets

- Loaded validated data onto Amazon Redshift using Copy Activity for end business consumption and analytics
- Automated SFTP egress to third-party vendors using Python scripts, eliminating 120+ hours of monthly manual work and reducing transmission errors by 90 percent

Accenture

Data Engineering Consultant

Gurugram, India

July 2021 – March 2023

- Architected real-time Spark streaming solution using Databricks and Apache Kafka with Informatica PowerExchange, processing 500,000+ events hourly with less than 60 second latency for near-real-time data availability
- Deployed streaming data pipelines with Slowly Changing Dimension (SCD-1) merge logic for PostgreSQL and BigQuery, synchronizing 8 production databases with 99.95 percent consistency through idempotent operations
- Implemented Delta Live Tables (DLT) framework for streaming ingestion with automated data quality enforcement and lineage tracking, reducing data quality incidents by 75 percent through declarative pipelines

Publicis Sapient

Associate Data Engineering L2

Gurugram, India

December 2019 – July 2021

- Led enterprise Cloudera-to-Databricks migration of 100+ YARN jobs, 80+ Apache Oozie workflows, and 12TB Azure Data Lake Storage (ADLS) data, completing 6 weeks ahead of schedule through parallel migration strategy
- Built Dataset Migration Framework using Python automation, migrating 200+ tables (partitioned and non-partitioned) from production Azure Synapse database to UAT environment in 8 hours versus 1 week manual effort, achieving 98 percent time reduction
- Designed push-down optimization framework for SQL Data Warehouse to ADLS extraction, boosting throughput by 50 percent and saving 15 hours weekly through predicate pushdown and parallel reads
- Designed comprehensive ETL workflows for Global Commercial Operations using Informatica PowerCenter, Informatica Cloud (IICS), and Big Data Management (BDM), integrating 12+ data sources serving 500+ business users

Tata Consultancy Services

Software Engineer - Data Engineering

Noida, India

November 2015 – December 2019

- Developed comprehensive ETL workflows for Global Commercial Operations using Informatica PowerCenter, Informatica Cloud, and Big Data Management, integrating 12+ data sources serving 500+ business users
- Automated GDPR compliance through personally identifiable information (PII) data masking scripts in Linux, eliminating 40 hours of fortnightly manual effort and ensuring 100 percent regulatory adherence across 10+ databases
- Built real-time Medical Inquiry data exchange solution integrating third-party REST APIs via Informatica Cloud, reducing response latency by 70 percent from 15 minutes to 4 minutes
- Resolved critical UTC-EST timezone discrepancies affecting 5 million+ records through custom transformation logic, ensuring accurate temporal data across 8 global time zones

EDUCATION

Bachelor of Engineering in Electrical and Electronics

Manipal University, Manipal, Karnataka, India

June 2011 – June 2015

CGPA: 8.30/10.0

CERTIFICATIONS

- **Databricks Certified Data Engineer Associate** – Validated expertise in Databricks platform, Delta Lake, Apache Spark, and data engineering workflows

AWARDS & RECOGNITION

- **Winner – AI & Data Center of Excellence Global Recognition 2025:** Partnering for Client Impact
- **Silver Cheers Awardee (December 2023):** Recognized for exceptional client impact and inclusive leadership
- **All Hands Star Award Nominee (March 2022, September 2022):** Honored for delivering measurable business outcomes and client value creation
- **Account Recognition Award (March 2020):** Celebrated for establishing UAT environment for Cloud Migration Project
- **On the Spot Awards (January 2019, February 2019):** Acknowledged for successful GDPR Automation and Informatica Cloud Migration implementation