



Mahesh Babu Kadiyala

Email: maheshbabukadiyala@gmail.com

Mobile: 8801234814

Professional Summary:

Data Engineering Manager with 14+ years of experience designing cloud-native batch and near real-time data platforms using Azure, Databricks, PySpark and Delta Lake. Experienced in Medallion Architecture, Structured Streaming, Azure Event Hubs, enterprise data governance and scalable Lakehouse solutions supporting analytics and AI workloads.

Core Skills & Expertise:

- ETL & Data Integration: ADF, SSIS, Informatica, Delta Lake, Snowflake
- Cloud & Data Platforms: Azure Data Lake, Azure Synapse Analytics, Azure Databricks, Azure Event Hubs (Kafka-compatible), Delta Lake
- Big Data & Processing: Python, PySpark, Spark SQL, Delta Lake, Data Warehousing
- Machine Learning & AI: Azure ML, AutoML, Predictive Analytics, LLM's
- Data Quality & Governance: Unity Catalogue, Schema Validation, ACID Compliance, Data Versioning
- DevOps & Performance Optimization: Azure DevOps (CI/CD), SQL Tuning, Query Optimization
- Security & Compliance: Data Governance Policies, Key Vault, Secure Data Pipelines
- Collaboration & Leadership: Cross-functional Team Management, Agile, Scrum

Certifications & Education:

- Microsoft Certified: Azure Data Engineer Associate (DP-203)
- Microsoft Certified: Data Analyst Associate (Power BI DA-100)
- ITIL Foundation V3 | Lean Six Sigma - Yellow Belt Certified
- PG Certification in Artificial Intelligence & Machine Learning – IIIT Hyderabad (Expected Jan 2026)
- Master of Computer Applications – JNTUK (2008-2011)

Work Summary:

Envoy Global Data Engg Manager / AI Solutions (Jul 2025 – Present)

- Leading the design and delivery of enterprise integration and AI solutions for global immigration and HR platforms using Azure, Databricks, Workato, REST APIs, and Azure OpenAI.
- Architected scalable integration solutions between Workday, HRIS platforms, Message Center, and enterprise applications using Workato recipes, secure REST APIs, webhook triggers, and event-driven workflows.
- Designed metadata-driven data processing frameworks enabling reusable, configurable pipelines across multiple enterprise clients, reducing development effort and improving maintainability.



- Developed AI-powered document processing and evaluation workflows leveraging Azure OpenAI (LLMs) to automate data extraction, validation, questionnaire evaluation, and workflow orchestration while maintaining human review checkpoints.
- Built secure API integrations using OAuth2, API Keys, Azure Key Vault, webhook authentication, retry mechanisms, and centralized logging to ensure reliable enterprise-grade integrations.
- Implemented reusable Workato connector frameworks supporting dynamic payload transformation, error handling, audit logging, and monitoring for business-critical HR and immigration workflows.
- Collaborated with product managers, architects, attorneys, and business stakeholders to translate complex business requirements into scalable cloud-native integration and AI solutions.
- Led technical design reviews, mentoring engineers on integration architecture, API design, security best practices, and AI adoption.

Alliantgroup, India Data Engineering Manager (Mar'23 – Jul'25sande)

- Led a globally distributed team of 6 data engineers across India and the USA, spearheading data engineering projects and delivering impactful solutions to drive business growth and innovation.
- Designed & implemented 50+ ETL pipelines using Azure Data Factory (ADF), PySpark, and Databricks, extracting millions of records from Salesforce and seamlessly loading them into Azure.
- Ensured 99.99% data integrity and security by implementing Azure Synapse Analytics and Data Lake solutions, adhering to industry-standard governance policies.
- Developed & optimized 100+ SQL procedures and scripts for data transformation, cleansing, and validation, reducing processing time by 30%.
- Improved query performance by 40% by fine-tuning data processing on Azure Synapse and Azure SQL, reducing report execution time from minutes to seconds.
- Automated CI/CD deployments using Azure DevOps, reducing manual intervention by 80% and accelerating release cycles.
- Implemented Unity Catalog governance for Databricks assets by organizing data into catalogs and schemas, enabling centralized access control, secure data sharing and improved governance across engineering teams.
- Led a cross-functional team of data engineers and analysts, ensuring zero downtime during complex data migrations.
- Designed & built a scalable Data Warehouse from the ground up using Azure Synapse Analytics, integrating 10+ data sources and engineering high-performance data models to support enterprise reporting.
- Unified structured and unstructured data by integrating Azure Data Factory, Data Lake, and SQL, enhancing business intelligence and improving decision-making speed by 35%.
- Developed Python scripts to process and consolidate 500+ legacy client files, transforming unstructured data into structured formats. This improved data accessibility and reporting efficiency by 50%, enabling faster insights.
- Designed and implemented Delta Live Tables (DLT) pipelines to automate Bronze, Silver, & Gold data processing, improving pipeline reliability through declarative transformations, built-in data quality expectations, automatic dependency management, and simplified orchestration.



- Automated Machine Learning (AutoML): Leveraged Azure AutoML to optimize model selection, hyperparameters, and feature engineering, reducing manual effort and accelerating data-driven innovation.

Centific, India

Lead Data Engineer

(Jul'22 – Feb'23)

- Architected and implemented a scalable ETL pipeline leveraging Azure Databricks to unify data from disparate source systems while adhering to the Medallion Architecture. The solution ingested raw data into the bronze layer, performed cleansing, deduplication, and enrichment in the silver layer, and transformed it into business-ready datasets in the gold layer for advanced analytics.
- Designed and implemented near real-time data ingestion pipelines using Azure Event Hubs and Databricks Structured Streaming to continuously ingest application events into Delta Lake, enabling low-latency analytics and downstream reporting. Implemented checkpointing, fault tolerance, and scalable micro-batch processing for reliable event delivery.
- Built streaming ETL pipelines using PySpark Structured Streaming to transform incoming event data, perform schema validation, deduplicate records, and write curated datasets into bronze, silver and gold delta tables.
- Processed customer activity events (login, transactions, order updates and application telemetry) through Azure Event Hubs, enabling near real-time operational dashboards and downstream analytics.
- Optimized Transformations: Implemented PySpark-based transformations within Databricks notebooks, enabling parallel processing for high-volume datasets, reducing processing time by 40%.
- Schema Evolution & Governance: Leveraged Delta Lake to enforce schema validation, maintain versioning, and ensure ACID compliance, improving data reliability.
- Performance Tuning: Optimized Spark cluster configurations, caching strategies, and partitioning techniques to enhance query performance and minimize latency.
- Business Intelligence & Reporting: Integrated refined Gold-layer data with Azure Synapse Analytics, enabling interactive Power BI dashboards and real-time reporting.

AkzoNobel, UK

Azure Data Engineer

(Nov'21 – Jun'22)

- Designed and automated Azure DevOps CI/CD pipelines, streamlining the deployment of 100+ database objects and Azure Analysis Services models, cutting deployment time by 40%.
- Partnered with stakeholders to define 10+ key performance indicators (KPIs), enabling data-driven decision-making.
- Led the design and implementation of scalable data models, optimizing storage and query performance.
- Boosted ETL efficiency, reducing Azure Data Factory pipeline execution times by 30% and optimizing data processing for millions of records.
- Developed multiple Power BI dashboards and 10+ tabular models, replacing manual Excel-based reports and reducing reporting time by 50%.
- Optimized data storage by integrating Azure BLOB and Data Lake Storage with Azure Synapse Analytics, managing terabytes of data.
- Architected end-to-end Azure Data Analytics solutions, enhancing data accessibility and business intelligence capabilities.



- Designed scalable and high-availability data models and architecture.
- Implemented Power Apps solutions for real-time data modifications, reducing manual intervention by 60%.
- Effectively communicated business goals, risks, and constraints to the delivery team, ensuring seamless execution and alignment with strategic objectives.

Eastman, India Advanced Systems Analyst (May'16 – Oct'21)

- Migrated on-prem SQL Server data to Azure Synapse & Snowflake, improving data accessibility.
- Implemented data validation and security controls, ensuring compliance with governance policies.
- Developed BI dashboards using Power BI, enhancing business insights and decision-making.

Previous Roles (Feb'13 – Apr'16)

- Lead Engineer, iGATE (Capgemini), India: Managed service continuity, optimized SQL queries, and automated ETL workflows.
- BI Consultant, I Tree, India: Designed data models, built SSRS reports, and optimized warehouse solutions.

Achievements & Impact:

- Designed and optimized 100+ ETL pipelines, ensuring seamless data flow across the organization.
- Improved performance of Snowflake & Databricks workloads, reducing processing time by 40%.
- Automated data validation and quality checks, achieving 99.99% accuracy.
- Led a globally distributed data engineering team, delivering enterprise-level data solutions.
- Integrated multiple cloud platforms (Azure, Snowflake, AWS), enabling scalable and flexible data architecture.