

PRASAD P R

+919740146523 | shenoy.prasad20@gmail.com | Bengaluru, Karnataka, India | [LinkedIn](#) | [GitHub](#) | [Website](#)

SUMMARY

Accomplished AI Engineer with 8+ years of expertise in designing, developing, and deploying both traditional AI/ML and LLM based solutions at scale. Proven success in optimizing LLMs, generative AI, and computer vision models for production, with deep proficiency in Python, PyTorch, TensorFlow, Langchain and OpenVINO. Experienced in building scalable pipelines, RAG pipelines, agentic workflows, model optimization frameworks, and AI-driven applications across NLP, RAG. Adept at leading AI projects end-to-end, mentoring engineers, and delivering solutions that drive measurable business impact.

SKILLS

Programming Languages: Python, JavaScript, C++

Machine Learning Techniques: Agentic AI, Generative AI, LLM, NLP, Deep Learning, Computer Vision

Data Processing and Analysis: Computer Vision, Image and Video Processing, Quantitative Analysis, Data Visualization Techniques, NoSQL

Software Development: Application Design, Microservices, Git, Docker, Kubernetes

Generative AI & LLMs: AI Agents, Multi Agents, MCP, Langgraph, CrewAi, AutoGen, LangChain, Hugging Face Transformers, OpenAI API, Azure OpenAI, Agentic AI, RAG (Retrieval-AugmentedGeneration), Prompt Engineering, Vector Databases (Pinecone, FAISS, Weaviate).

Frameworks & Libraries: Pytorch, TensorFlow, OpenVINO, ONNX.

MLOps & Deployment: Fast API for AI microservices, Docker, Kubernetes, Django REST Framework

Databases: PGSQL, MongoDB

PROFESSIONAL EXPERIENCE

Senior Associate, PWC India

Oct '25 - Present

- Specialized in LangGraph, multi-agent systems, and agentic workflows for autonomous AI architectures with complex state management and orchestration.
- Enhanced proficiency in production-ready AI agent design, autonomous decision-making systems, and enterprise-scale multi-agent coordination.

Cloud Application Development Engineer, Intel Corporation

Aug '18 - Sept '25

- Achieved **10x performance gains** on Intel GPU by optimizing Hugging Face transformer models (Whisper, M2M100) using ONNX Runtime and OpenVINO; designed automated benchmarking pipeline reducing evaluation time by **90%** (5 hours → 30 minutes).
- Engineered real-time 3D camera calibration and point cloud segmentation pipeline using PyTorch, Open3D, and Intel RealSense, achieving **80% reduction in inference latency** and 50% improvement in processing efficiency.
- Optimized AI models for **30% faster training/inference** and deployed to IoT-enabled edge devices with MQTT protocols and real-time video streaming capabilities.
- Led development of MERN stack Solution Delivery System serving **1,000+ users**, architecting microservices migration (30% faster access); spearheaded CI/CD integration reducing maintenance overhead by **80%**.

Software Engineering Intern, Intel Corporation

Oct '17 - Aug '18

- Contributed code to 2 repositories in the GitHub Archives program as a part of Hyperledger Sawtooth Enterprise Blockchain.
- Built AngularJS UI for blockchain visualization and monitoring
- Developed POCs for Stereo Vision, Computer Vision, and AI-based node analytics.

EDUCATION

Liverpool John Moores University - Liverpool & upGrad

Apr '25 -- Present

M.Sc. in Artificial Intelligence and Machine Learning

IIIT-B & upGrad

Feb '24 — Apr '25

Executive PG Diploma in Artificial Intelligence and Machine Learning

3.63/4 CGPA

JSS Academy of Technical Education, Bangalore

Aug '14 — Jul '18

Bachelor of Engineering (B.E.), Computer Science & Engineering

Grade – Distinction

KEY PROJECTS (Above and Beyond my Work)

Find Me A Cube (Internal tool)

Objective: No cost solution for identifying empty cubes in a Hybrid Model. Also provide cube usage analytics.

Solution: Cloud hosted solution using Django REST API as a backend, React JS based frontend and Mongo DB

database. **Result:** Faster, hassle free cube identification for employees, the clean frontend and faster backend.

Automated Smart Reports along with Enterprise RAG (Internal Tool)

Objective: Speed up report collation, reduce manual report generation based on multiple mails. **Solution:** Using Gen

AI & prompt engineering, extract required information from multiple reports and combine the information into a

single status report. **Result:** Reduced the managers time spent on report generation from 5 days to 30 minutes.

Smart Hiring and Resume Screening using Generative AI (Internal Tool)

Objective: Help hiring managers screen suitable resumes based on the job description required for the role. **Solution:**

Using LLM models, understanding the data of the resume and then trying to perform a similarity mapping with the job

description to obtain the top resumes that are suitable for the job. **Result:** Helped managers reduce the resume

screening time from hours to minutes.

AWARDS & CERTIFICATIONS

Divisional Recognition Award (DRA)

Intel

Innovation Star

Intel

Deploy an Agent with Google Agent Development Kit (ADK)

[Google](#)

IBM RAG and Agentic AI Professional Certificate

[Coursera](#)

LangChain- Develop LLM powered applications with LangChain

[Udemy](#)

Generative AI with Large Language Models

[Coursera](#)