

RAJESH SANTHA

Senior Data Engineer · Spark | PySpark | Databricks | Airflow | Azure

+91 70322 16316 | rajesh.santha91@gmail.com | Hyderabad, India | B.Tech – Information Technology – JNTU Hyderabad

CAREER SUMMARY

- Senior Data Engineer with 11+ years of experience building scalable batch and real-time data platforms using Spark (Scala/PySpark), Azure Databricks, Airflow, Delta Lake, Kafka, and Azure cloud technologies.
- Proven expertise in designing reusable ingestion frameworks, optimizing distributed Spark pipelines, and delivering high-performance cloud-based data solutions across banking, aviation, telecom, retail, and insurance domains.
- Azure & Databricks certified professional with a strong foundation in computer science and distributed systems.

TECHNICAL SKILLS

Languages: Scala, Python

Big Data & Streaming: Apache Spark, PySpark, Spark SQL, Spark Structured Streaming, Delta Lake, Apache Kafka, Hive, Cloudera CDP

Cloud & Platforms: Azure Databricks, ADLS Gen2, Azure Data Factory (ADF), Kubernetes (AKS), Docker, Snowflake, SQL Server, MongoDB, Elasticsearch, HBase, HDFS

Architecture: Medallion Architecture (Bronze / Silver / Gold), Lambda & Kappa patterns

Orchestration: Apache Airflow, Azure Data Factory, Control-M, Apache Oozie

CI/CD & Version Control: Git, GitLab CI/CD, Azure DevOps, Bitbucket, Azure Repos, Maven, SBT

Tools: IntelliJ IDEA, Eclipse, Jira, Confluence, ServiceNow

CERTIFICATIONS

- Databricks Certified Data Engineer Associate
- Azure Data Engineer Associate (DP-203)
- Azure Data Scientist Associate (DP-100)
- Azure Data Fundamentals (DP-900)
- Azure Fundamentals (AZ-900)
- Azure AI Fundamentals (AI-900)
- Cloudera Certified Spark & Hadoop Developer (CCA175)
- Oracle Certified Java SE 8 Programmer I (OCA)
- MuleSoft Certified Integration & API Developer (MCD Associate)

PROFESSIONAL EXPERIENCE

UBS · Associate Director – Data Engineering · Hyderabad

May 2025 – Present

Coforge · Technical Analyst · Hyderabad

Sep 2022 – May 2025

Mindtree · Module Lead · Bangalore

Jul 2020 – Aug 2022

Tech Mahindra · Senior Data Engineer · Bangalore

Apr 2019 – Jul 2020

Accenture · Software Engineer · Hyderabad

Apr 2015 – Apr 2019

PROJECT EXPERIENCE

Excalibur Data Framework – UBS — Associate Director – Data Engineering

May 2025 – Present

Tech: Spark (Scala/PySpark), Azure Databricks, Delta Lake, ADLS Gen2, PostgreSQL, GitLab CI/CD, Kubernetes (AKS), Apache Airflow

- Designed and developed scalable Spark-based ingestion frameworks to process high-volume market data into Delta Lake on Azure Databricks.
- Built curated ADLS-to-PostgreSQL publishing pipelines for downstream analytics and reporting applications.
- Integrated Airflow for orchestration and scheduling of dependent Spark jobs with SLA monitoring, retry handling, and failure recovery.
- Optimized distributed Spark pipelines via partitioning, caching, and join tuning for large-scale financial data workloads.
- Implemented CI/CD pipelines using GitLab for automated build, testing, version control, and deployment of Spark and PySpark applications.

Emirates Airlines – Coforge — Technical Analyst

Sep 2022 – May 2025

Tech: Spark (Scala), PySpark, Azure Databricks, Snowflake, Cloudera CDP, Apache Oozie, Azure Data Factory (ADF)

- Developed enterprise data ingestion pipelines across raw, branded, LMO, and modelling layers using Spark (Scala) and Azure Databricks.
- Designed and maintained Oozie workflows for orchestrating large-scale batch pipelines and downstream Snowflake integrations.
- Built external tables and views in Snowflake to support business reporting and analytics use cases.
- Developed API-based ingestion frameworks to extract, flatten, and process large-scale JSON data into curated Snowflake tables.
- Contributed to migration of enterprise data pipelines from Cloudera CDP to Azure, replacing legacy Oozie workflows with Azure Data Factory.

BDL Harmonization – Unilever (via Mindtree) — Senior Azure Data Engineer

Dec 2020 – Aug 2022

Tech: Spark (Scala), PySpark, Azure Databricks, Azure Data Factory (ADF), ADLS Gen2, Azure DevOps, SBT

- Built scalable Azure-based data pipelines for ingesting and processing enterprise data using Spark, PySpark, and Azure Databricks.
- Ingested data in multiple formats (Parquet, JSON, CSV) using Spark on Azure Databricks with Spark table registration in ADLS Gen2.
- Developed orchestration pipelines in Azure Data Factory for full and incremental data loads.
- Implemented complex Spark transformation logic applying optimization techniques including partitioning, caching, and broadcast joins.
- Part of the architecture team for defining data structures; collaborated with client-side data engineering team on schema design.
- Collaborated with Product Owners, Architects, and Business Analysts to design end-to-end enterprise data flows.
- Implemented CI/CD pipelines via Azure DevOps for automated deployment and release management.

MLS Data Pipeline – Realogy (via Mindtree) — Senior Data Engineer

Jul 2020 – Nov 2020

Tech: Apache Spark, Scala, Apache Kafka, Apache NiFi, Java Microservices, MongoDB, PostgreSQL, AWS Batch, AWS S3

- Developed distributed Spark pipelines for processing MLS listings, photos, and large-scale real estate datasets.
- Built Spark UDFs enabling business users to perform custom transformations through UI-driven workflows.
- Implemented end-to-end tracking and monitoring of MLS data pipelines using Apache NiFi and Apache Kafka.

Telecom Event Management – T-Mobile (via Tech Mahindra) — Senior Data Engineer

Apr 2019 – Jul 2020

Tech: Spark Structured Streaming, Scala, Apache Kafka, Elasticsearch, Kibana, SQL

- Developed real-time SNMP event management and Kafka-based streaming pipelines using Spark Structured Streaming, Scala, and Elasticsearch.
- Implemented DataStream joins, dynamic mapping refresh mechanisms, and stream transformation logic for large-scale network alarm processing.
- Reduced alarm processing traffic by 25% through threshold delay implementation, event deduplication, and reusable ingestion framework development.

Anthem DataLake – Anthem Inc. (via Accenture) — Software Engineer

Apr 2015 – Apr 2019

Tech: HDFS, Apache Spark, Spark Streaming, Scala, Apache Kafka, Hive, HBase, Sqoop, StreamSets, Control-M

- Performed large-scale insurance data analysis and distributed processing using Hadoop, Apache Spark, Scala, and Kafka.
- Developed enterprise data ingestion and Hive-to-HBase publishing frameworks processing data from Oracle, DB2, MySQL, JSON, CSV, TEXT, and EBCDIC sources.
- Built automation and workflow scheduling frameworks using Spark, Scala, StreamSets, and Control-M for enterprise batch data processing.