

Sourav Suman

Gen AI Developer

✉ sourav2fly@gmail.com

☎ 8059944096

📍 Bengaluru, India

🔗 tinyurl.com/2psz5fu8

🌐 linkedin.com/in/srvsmn/

🐙 github.com/srv-smn

Education

BE: Computer Science and Engineering,

SSGMCE, Maharashtra, India
Maharashtra, India

Skills

Languages

- Python
- JS
- Solidity

Technology

- Model Finetuning, Agents
- Workflow
- Voice Agents
- RAG
- Knowledge Graph
- Transformer Architecture
- Blockchain

Gen AI Frameworks

- Langchain
- Langgraph
- AutoGen
- LiveKit

Backend

- Node.js
- Fast API
- Subgraph

Profile

A Lead Developer with 5 years of experience, now specializing in building practical applications with Generative AI. Based in Bangalore, India, my work involves architecting complex, agent-based systems for real-time voice and data processing using frameworks like LangChain, LangGraph, and AutoGen. I also specialize in fine-tuning Large Language Models (LLMs) and building scalable, privacy-aware solutions by implementing advanced RAG pipelines and Knowledge Graphs.

Professional Experience

Lead Developer, knuiK Innovnition

08/2022 – Present

- Led end-to-end development of Generative AI solutions, including intelligent agents, RAG pipelines, and fine-tuned LLMs for real-world applications.
- Designed multi-step AI workflows and integrated GenAI capabilities into products using tool use, MCP's.
- Worked on project estimation, system architecture, and development of blockchain applications, including smart contracts and Subgraph integrations on EVM chains.

Software Developer, Fcusisicni S;sicms, Func

07/2021 – 08/2022 | Pune

- Developed smart contracts for various blockchain platforms to meet project requirements.
- Integrated smart contracts with existing systems to enhance functionality and efficiency.
- Conducted unit testing and implemented best practices to ensure code quality.
- Collaborated with cross-functional teams to ensure successful project delivery.

Projects

PMMIC: Vision-Language Model for Industrial

Compliance ☑

07/2025

- Trained specialized vision models (like Qwen and Gemini) to automatically find safety and compliance problems in photos of petroleum tanks.
- Used efficient fine-tuning methods (SFT with LoRA) and reinforcement learning (GRPO) to build a highly accurate, low-cost model.
- The final AI is skilled at spotting both common and rare edge-case issues, making inspections faster and more reliable.

Technologies: Qwen2VL, Gemini 1.5 Flash, LoRA, SFT, GRPO, LLAMA Factory

Hachathons

Google Agentic AI Day - Finalist

07/2025

Got select to build the solution based on AI that can help farmers.

Edu Chain Sem 3 Winner

06/2025

Developed a AI powered gamified learning platform with OCID on edu chain.

Outbound SIP Call Agent: High-Volume Conversational AI 01/2025

- Built a powerful AI agent that makes hundreds of human-like phone calls at once, perfect for automating large-scale customer outreach or lead qualification.
- Kept conversations feeling natural and real-time by engineering a low-latency (~800ms) system using LiveKit for fast audio processing.
- Designed a flexible system where different AI services (speech-to-text, language models) can be easily swapped, and used a RAG pipeline to give the AI a knowledge base for smarter responses.

Technologies: Python, LiveKit, Twilio (SIP), RAG, AI Workflows, LLMs, STT/TTS

Hatchala: Automated Tax Filing with AI h Privacy-Preserving Workflows

04/2024

- Built an AI-driven system to classify tax forms, extract key data, and auto-file with minimal human input.
- Ensured privacy by masking sensitive data before LLM use and deploying secure, agent-based workflows.
- Combined **OCR + LLMs** for accurate extraction and used **RPA** for seamless end-to-end filing.

Technologies: Python, OpenAI, Gemini, Qwen, LangGraph, Tesseract, Custom AI Agents, RPA

MyHoudini: AI Voice Companion for Elderly Support 08/2024

- Developed a friendly, voice-controlled AI assistant to help elderly individuals manage daily tasks like reminders and scheduling, while also providing companionship.
- Designed a cost-effective system using LiveKit for clear audio and a smart pipeline for the AI's speech processing, making the technology affordable and accessible.
- The assistant works hands-free, allowing users to simply talk to it to store memories, set appointments, or just have a conversation, making it easy and natural to use.

Technologies: LiveKit, Python, Large Language Models (LLMs), AI Agents, Speech-to-Text (STT), Text-to-Speech (TTS)